

RLChina Reinforcement Learning Summer School



RLChina 2022

PyTorch Introduction

YanSong



CASIA


August 15, 2022

Outline

- Why Pytorch?
- Introduction to Pytorch tensor operations
- BackPropogation
- Training
- Examples and Homeworks
- Jidi platform

 PyTorch



及第 


What is PyTorch?

- Replacement for **Numpy** to use the power of **GPUs**
- **Deep learning research** platform that provides flexibility and speed
- Installation: <https://pytorch.org/get-started/locally/>
- Many alternatives based on similar principles




Why PyTorch?

- **Dynamic computation graphs** (by now also supported in **TF Eager**) are from our experience more intuitive for newcomers
- Easy debugging
- Data Parallelism




Scalars, Vectors, Matrices, Tensors and Basic operations

$$\vec{X} = \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

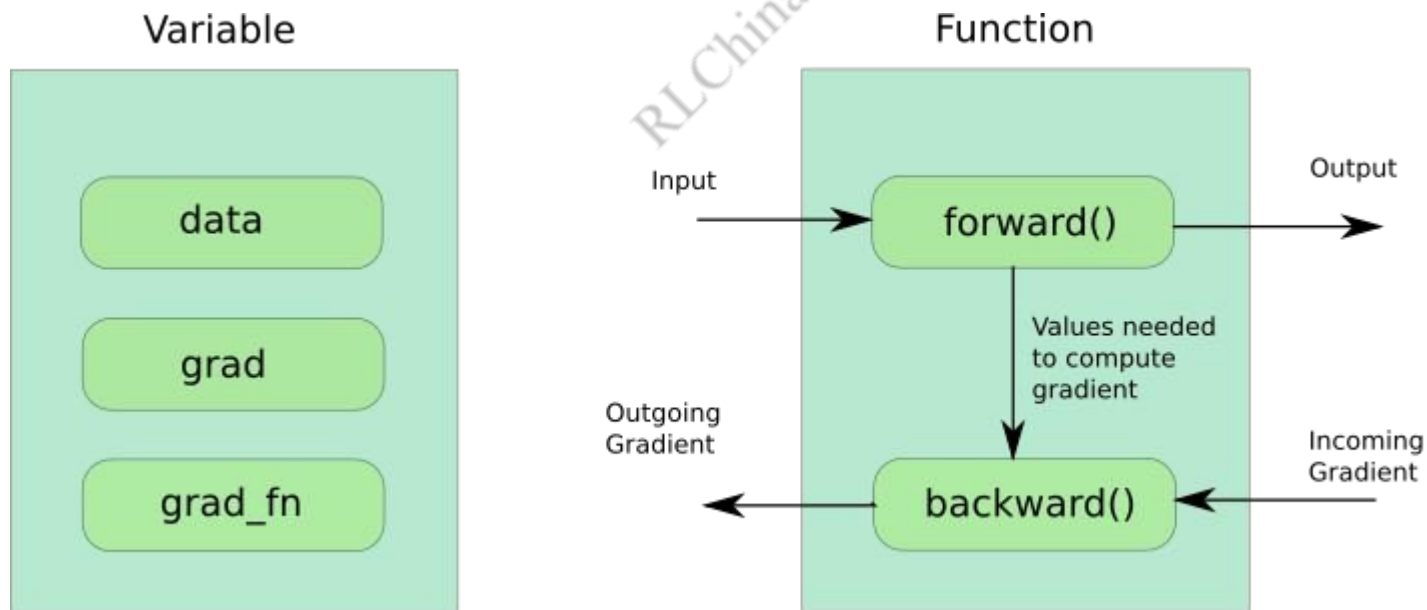


$$\mathbf{X} = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1n} \\ x_{21} & \ddots & & x_{2n} \\ \vdots & & \ddots & \vdots \\ x_{m1} & x_{m2} & \cdots & x_{mn} \end{bmatrix}$$



PyTorch Variables

- A PyTorch Variable is a **wrapper** around a PyTorch Tensor, and represents a **node** in a computational graph.
- If x is a Variable then $x.data$ is a **Tensor** giving its value
- $x.grad$ is another Variable holding the **gradient** of x with respect to some scalar value

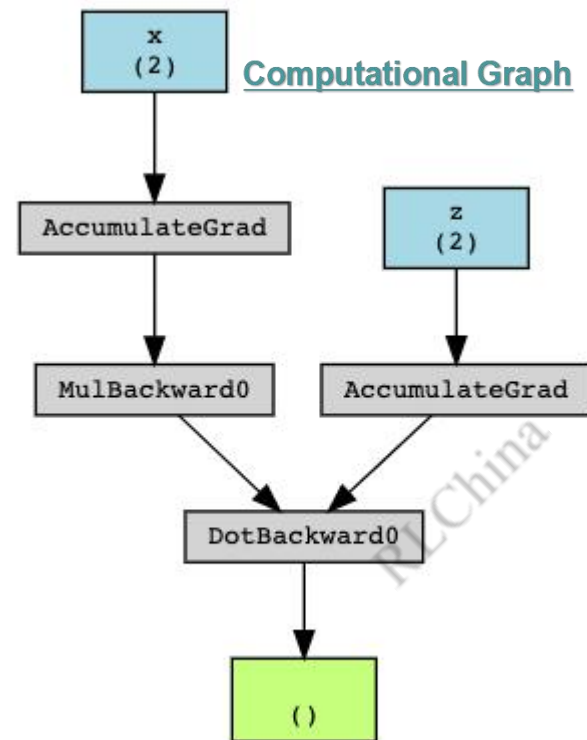


PyTorch Autograd---Gradient

```
x = torch.tensor([-1.5, 1.2], requires_grad=True)
y = torch.tensor([1.0, -1.3])
z = torch.tensor([-2.0, 0.2], requires_grad=True)
r = x * y @ z
r.backward()
x.grad
> tensor([-2.0000, -0.2600])
```

$$\begin{aligned}\frac{\partial}{\partial \mathbf{x}} r &= \frac{\partial}{\partial \mathbf{x}} (\mathbf{x} \odot \mathbf{y})^\top \mathbf{z} \\ &= \mathbf{z} \left[\frac{\partial}{\partial \mathbf{x}} (\mathbf{x} \odot \mathbf{y}) \right] + (\mathbf{x} \odot \mathbf{y}) \left[\frac{\partial}{\partial \mathbf{x}} \mathbf{z} \right] \\ &= \mathbf{z} \left[\mathbf{y} \left[\frac{\partial}{\partial \mathbf{x}} \mathbf{x} \right] + \mathbf{x} \left[\frac{\partial}{\partial \mathbf{x}} \mathbf{y} \right] \right] \\ &= \mathbf{z} \odot \mathbf{y} = \begin{bmatrix} -2.0 \\ 0.2 \end{bmatrix} \odot \begin{bmatrix} 1.0 \\ -1.3 \end{bmatrix} = \begin{bmatrix} -2.00 \\ -0.26 \end{bmatrix}\end{aligned}$$

BackPropagation derivation (Chain Rule)



$$\mathbf{x}_{t+1} = \mathbf{x}_t - \alpha \frac{\partial}{\partial \mathbf{x}_t} r_t$$

$$\theta_{t+1} = \theta_t - \alpha \frac{\partial}{\partial \theta_t} L(X, Y; \theta_t)$$

Gradient update

code on supplementary notebook

PyTorch Autograd---BackPropagation

- Backpropagation = Efficient Application of Chain Rule

- Chain Rule:
$$y = g(x)$$
$$z = f(y)$$
$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = f'(g(x))g'(x)$$
- Backprop



Naomi Saphra
@nsaphra

Following

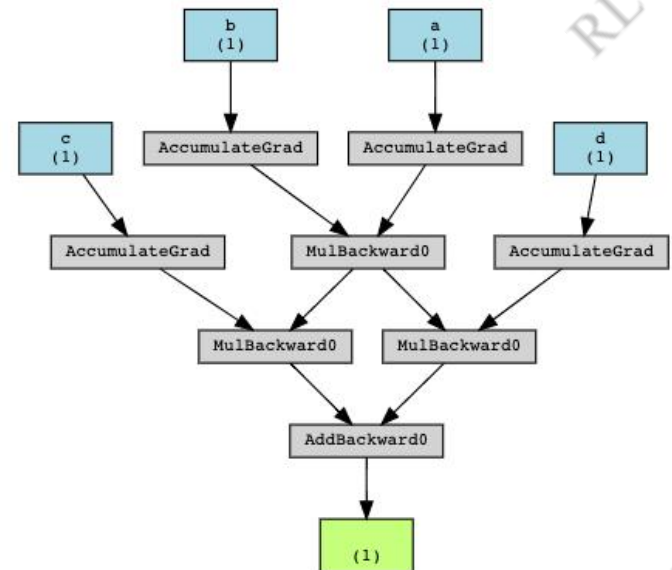
What idiot called it "deep learning hype" and not "backpropaganda"

3:05 PM - 14 Apr 2016

```
a = torch.rand(1, requires_grad=True)
b = torch.rand(1, requires_grad=True)
c = torch.rand(1, requires_grad=True)
d = torch.rand(1, requires_grad=True)
```

```
e = a * b
# compare to: f = c * a * b + d * a * b
f = c * e + d * e
```

```
f
> tensor([0.0994], grad_fn=<AddBackward0>)
```



code on supplementary notebook

PyTorch Modules

- All network components should inherit from `nn.Module` and override the `forward` method
- Using a module provides functionality:
 - Keeps track of trainable parameters
 - Allows you to easily swap between CPU and GPU (see `.to(device)`)
- To register a variable tensor to the parameters of a module you need to wrap it using `nn.Parameter`

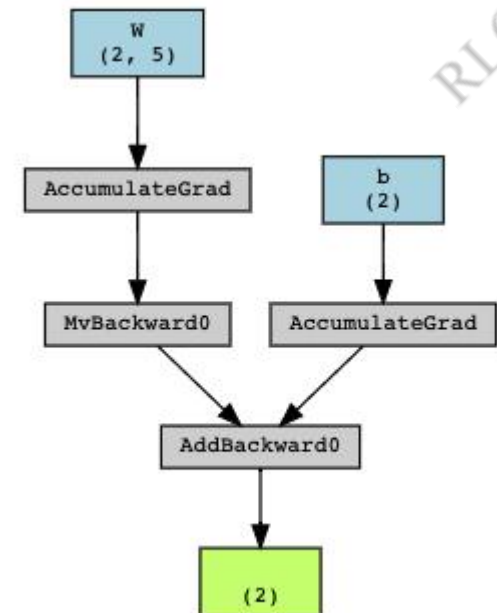
```
class LinearModule(torch.nn.Module):
    def __init__(self, x_dim, y_dim):
        super(LinearModule, self).__init__()
        self.W = nn.Parameter(torch.randn(y_dim, x_dim, requires_grad=True))
        self.b = nn.Parameter(torch.randn(y_dim), requires_grad=True)
    def forward(self, x):
        return self.W @ x + self.b

# Some random input and output data
x = torch.randn(5)
y = torch.randn(2)

model = LinearModule(5, 2)

for param in model.parameters():
    print(param.size())

pred = model(x)
pred
```



Loss Function


- Least squares [`nn.MSELoss`] $L(f_{\theta}(x), y) = \frac{1}{2} (f_{\theta}(x) - y)^2$
- Logistic [`nn.SoftMarginLoss`] $L(f_{\theta}(x), y) = \log(1 + \exp(-yf_{\theta}(x)))$
- Hinge loss [`nn.MultiMarginLoss` / `nn.MultiLabelMarginLoss`]
 $L(f_{\theta}(x), y) = \max(0, 1 - yf_{\theta}(x))$
- Cross-entropy [`nn.CrossEntropyLoss`]
 $L(f_{\theta}(x), y) = - [y \log(f_{\theta}(x)) - (1 - y) \log(1 - f_{\theta}(x))]$

For more see <https://pytorch.org/docs/stable/nn.html#loss-functions>

Training Loop


- Given model f_{θ} , initialize parameters θ (e.g. randomly)
- For number of epochs:
 - for number of iterations (i.e. number of batches in data):
 - sample batch of data $\mathbf{x}, \mathbf{y} \sim \mathbf{T}$
 - \mathbf{x} : input, \mathbf{y} : target output
 - run model forward $\mathbf{y}^* = f_{\theta}(\mathbf{x})$ and calculate loss $L(\mathbf{y}^*, \mathbf{y})$
 - Calculate gradient of loss $\nabla_{\theta} L$ w.r.t parameter using backprop
 - Update parameters using optimiser, e.g. $\theta_{t+1} = \theta_t - \alpha \nabla_{\theta} L$

Stochastics Gradient Descent



Li, Hao, et al. "Visualizing the loss landscape of neural nets." *NeurIPS*. 2018.

Optimiser



- `torch.optim.SGD`
- `torch.optim.Adagrad`
- `torch.optim.Adadelata`
- `torch.optim.Adam`
- `torch.optim.RMSprop`
- ... and many more

PyTorch Training Loop Scaffold

```
import torch.nn as nn
import torch.optim as optim

# Set a seed; your experiments should be reproducible!
torch.manual_seed(1)
# Load data
train, dev, test = ...
# Instantiate model
model = MyModel(...)
# Define loss function
loss_fn = nn.CrossEntropyLoss()
# Instantiate optimizer with a learning rate (lr)
optimizer = optim.SGD(model.parameters(), lr=0.1)

for epoch in range(10): # 10 epochs in this example
    for i, batch in enumerate(train): # assuming train is a generator that reshuffles data
        # Set gradients to zero
        optimizer.zero_grad()
        # Run forward
        y = model(batch)
        # Calculate loss
        loss = loss_fn(y, y_target)
        # Run backward to compute gradient of loss w.r.t. to model parameters
        loss.backward()
        # Perform one step of optimization
        optimizer.step()
        # Print diagnostics (e.g. loss or dev set performance)
        ...

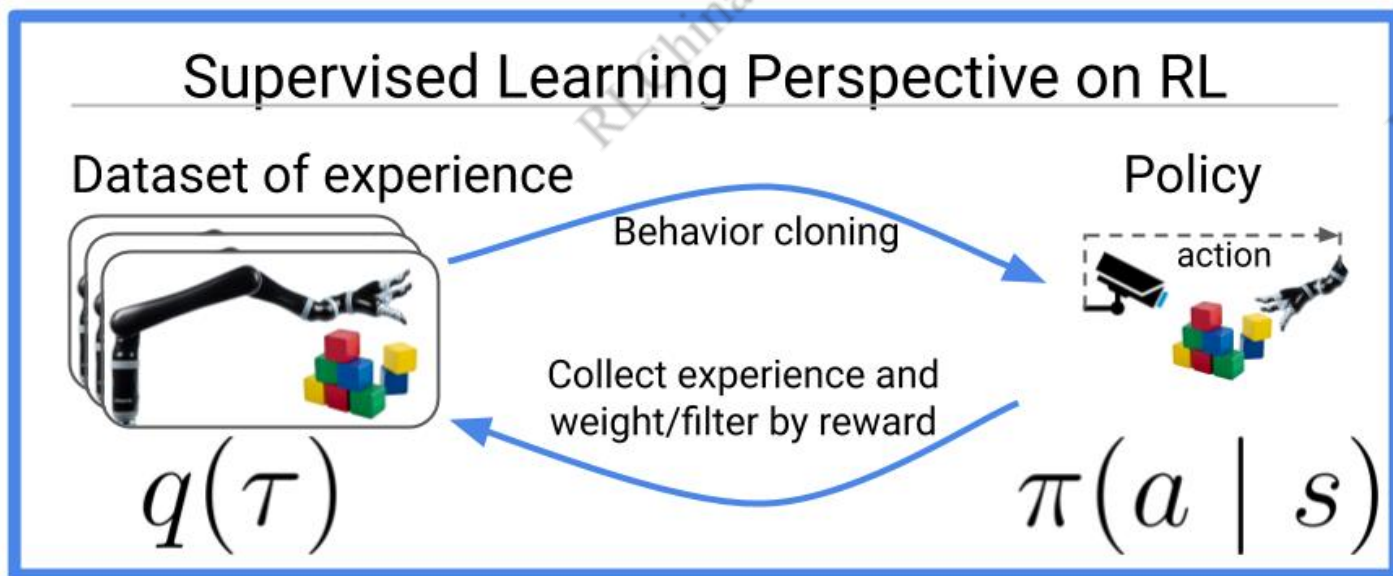
# Evaluate on test
...
```

Examples

- Regression
- Classification
- Recurrent Neural Network
- Convolution Neural Network



A Decision Making example using PyTorch

- **Behaviour Cloning** is a supervised learning way to tackle a decision making problem.
- Given expert experience, our goal is to imitate its behaviour as similar as possible with a parametrised policy
- therefore a good example to work with before get into the field of RL.



Homework: Behaviour Cloning on Olympics Wrestling

- two agent compete on stage and aim at pushing opponent out of bounds while avoid themselves falling out of the stage.
- Expert policy knows how to stay on the stage to survive.
- You are given with 10,000 collected expert data, battle begins!!!



How to submit your policy ---- Jidi



1. Align your policy input and output with Jidi's evaluation, see `run_log.py` in `ai_lib` (https://github.com/jidiai/ai_lib) and `/submission_example` folder.
2. Create an account on Jidi (<http://www.jidiai.cn/>)
3. Submit your policy



Thanks